

Fys4150 Project 1

Aksel Graneng

September 11, 2019

Abstract

This paper looks at solving differential equations as sets of linear equations, and concludes with the importance of writing non-generalized code for matrix operations.

1 Introduction

In this project, we will be solving a differential equation as a set of linear equations. This will be done using algorithms of varying generalization in python.

We will look at just how row reduction can solve a differential equation, as well as the efficiency in different row reduction methods.

2 Theory

2.1 Vectorized second derivative

If we have a 1d data-set on the form:

$$V(\vec{x}) = [v_0, v_1, \dots, v_{n-1}, v_n, v_{n+1}, \dots]$$

Then we can write the second derivative of the data-set as:

$$f_n = -\frac{v_{n+1} + v_{n-1} - 2v_n}{\Delta x^2}$$

Where Δx is the change in variable we are derivating based on; usually time.

Rather than calculating the second derivatives of this data-set individually,

we can instead calculate them all at the same time using linear algebra. This can be done by finding a matrix A such that

$$A\vec{V} = \vec{f}$$

Where:

$$\vec{f} = \begin{bmatrix} 2v_i - v_2 \\ -v_1 + 2v_2 - v_3 \\ \vdots \\ -v_{n-1} + 2v_n - v_{n+1} \\ -v_n + 2v_{n+1} - v_{n+2} \\ \vdots \end{bmatrix}$$

As multiplying \vec{f} with $\frac{1}{\Delta x^2}$ would give us an array containing all the second derivatives. We can see that \mathbf{A} must be:

$$\mathbf{A} = \begin{bmatrix} 2 & -1 & 0 & 0 & \cdots \\ -1 & 2 & -1 & 0 & \cdots \\ 0 & -1 & 2 & -1 & \cdots \\ \vdots & 0 & -1 & 2 & -1 \end{bmatrix}$$

This matrix can then be used to put up a set of linear equations:

$$\mathbf{A}\mathbf{v} = \mathbf{f}$$

Where \mathbf{v} contains solutions to the differential equation:

$$-u''(x) = f(x)$$

3 Method

3.1 General linear equation solver

The first thing our program does is solve a set of general linear equations on the form $\mathbf{A}\mathbf{v} = \mathbf{f}$. This is done by putting the matrix \mathbf{M} on row reduced echelon form:

$$\mathbf{M} = \begin{bmatrix} b_1 & c_1 & 0 & 0 & \cdots f_1 \\ a_2 & b_2 & c_2 & 0 & \cdots f_2 \\ 0 & a_3 & b_3 & b_3 & \cdots f_3 \\ \vdots & 0 & a_4 & b_4 & \cdots f_4 \end{bmatrix}$$

Where \mathbf{M} is a $n \times (n + 1)$ matrix. This is generally quite simple to solve using forward and backward substitution, but when n becomes very large (in my experience, larger than 10^5) we get problems with memory as well as computation time, as the most simple row reduction algorithm has $2(n + 1)n^2$ FLOPS.

To avoid this problem, we use the fact that most of the elements in the matrix is zero. In fact, there are at most 4 elements in each row that isnt zero. By removing every element that is zero from the matrix (except one in for the first and last line), we get:

$$\mathbf{M}^* = \begin{bmatrix} b_1 & c_1 & 0 & f_1 \\ a_2 & b_2 & c_2 & f_2 \\ \vdots & \vdots & \vdots & \vdots \\ 0 & a_n & b_n & f_n \end{bmatrix}$$

So now we have a $n \times 4$ matrix which contains all the information we need.

We now use forward substitution on \mathbf{M}^* just like we would on \mathbf{M} , except this time we "roll" the next row every time we subtract, so that b_i is subtracted from a_{i+1} and c_i is subtracted from b_{i+1} . At the same time, we are careful to subtract f_i from f_{i+1} . This means that index notation must be used, rather than `numpy.roll`.

This is done in a loop that has n iterations. First it does 3 operations, one on each element in the row except for the first (which is zero), along with a division, to normalize it. Then it does 3 subtraction and multiplication operations on the next row.

Next is backwards substitution. Again we use a loop with n iterations. Here we only do 2 subtractions and multiplications, on the last and second last element of each row.

All in all, this sums up to $14n$ FLOPS.

After it was done, the general linear equation solver was tested on:

$$u''(x) = f(x)$$

$$f(x) = 100e^{-10x}$$

Which has a closed-form solution:

$$u(x) = 1 - (1 - e^{-10})x - e^{-10x}$$

3.2 Specific linear equation solver

The matrix we are working with in this project is very simple. All values of a , all values of b and all values of c are equal. That is to say:

$$\text{for any } i \wedge j \in [1, n], \quad a_i = a_j, \quad b_i = b_j, \quad c_i = c_j$$

This means that we are practically doing the same operations n times when we are row reducing our matrix. Rather than doing this, we can limit our algorithm to only operate on our values for f , as these are unknown. We simply need to look at what the row reduction algorithm does to the matrix to see what it does to f , and then skip right to that part when calculating it numerically.

Studying the forward substitution, we see, starting with $n = 2$:

$$f_n^* = \frac{n}{n+1}(f_n - f_{n-1}^*)$$

Thus, the n 'th row ends up as (excluding all the zeros):

$$[0, 1, -\frac{n}{n+1}, f_n^*]$$

The very last row ends up as (with $n = \text{size of } f$):

$$[0, 0, 1, -(f_{-1} - f_{-2}^*)\frac{n}{n+1}]$$

The first value of f ($n = 1$) ends up as:

$$f_1^* = \frac{1}{2}f_1$$

And then we look at backwards substitution. We can see that we subtract:

$$\frac{n-i}{n-i-1}f_{-(i-1)}^{**}$$

from f_{-i}^* to find our final value f_{-i}^{**}

The process of forward substitution ends with the last line being:

$$f_1^* = f_1^* + 2f_2^{**}$$

The specified algorithm for solving the linear equations is a lot more efficient than the general one. Assuming we have already calculated the required numbers used in finding f , the forwards and backwards substitutions each require 2 operations, resulting in a total of $4n$ FLOPS.

3.3 LU Decomposition, `numpy.linalg.solve`

We should be comparing our results to what we get from using a built-in LU-decomposition method. This seems to require some work in python, so instead of using `numpy.linalg.lu` and doing calculations, we have opted to use `numpy.linalg.solve`, which fully solves the set of linear equations.

4 Results

4.1 General linear equation solver

We tested the general linear equation solver for datapoints $n = 10$, $n = 100$, $n = 1000$

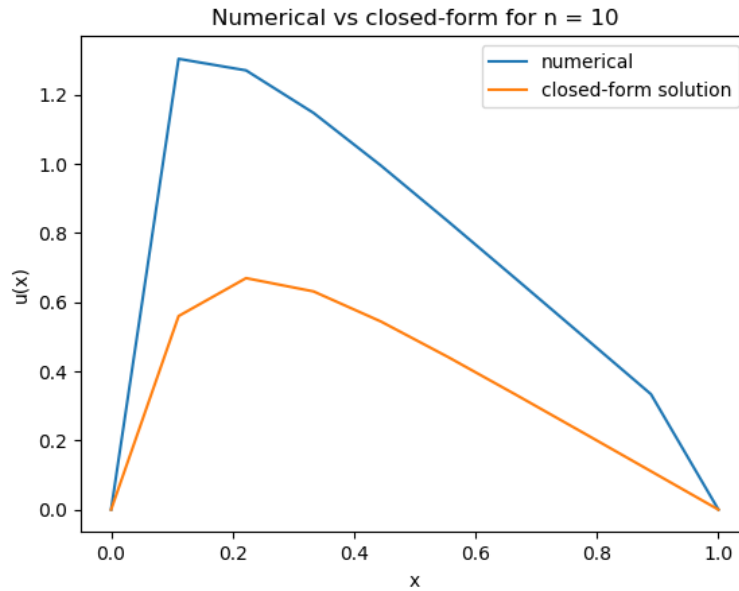


Figure 1: The numerical solution vs the closed-form solution for 10 data-points. Here we can see that they deviate massively.




Figure 2: The numerical solution vs the closed-form solution for 100 data-points. Here we can see that the difference is decreasing.

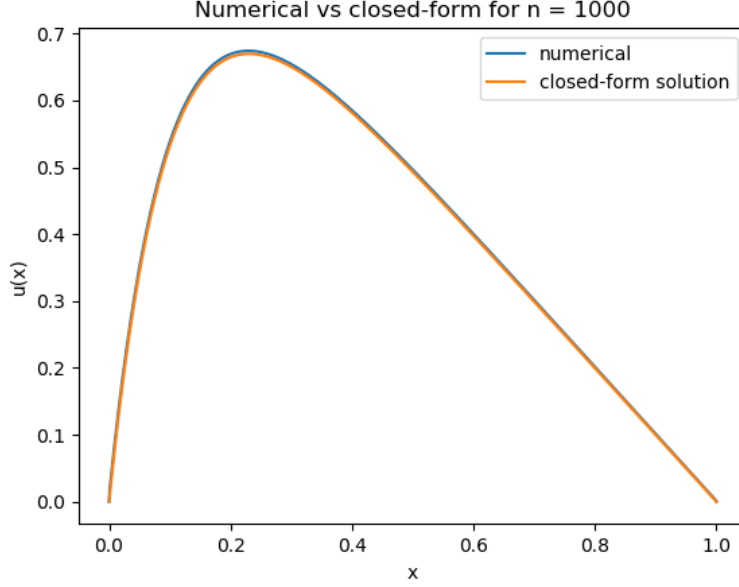


Figure 3: The numerical solution vs the closed-form solution for 100 data-points. Here we can see that the difference has been even more decreased, and it is becoming hard to tell them apart.

4.2 Computation time test

The general linear equation solver method took 6.4 seconds to compute 10^6 data points.

The specific linear equation solver method took 1.0 seconds to compute 10^6 data points.

4.3 Relative error

Using the specific linear equation solver method for data points ranging from 10 to 10^7 , we got these relative errors:

n	10^1	10^2	10^3	10^4	10^5	10^6	10^7
ϵ_{\max}	10^{-1}	$5.3 \cdot 10^{-3}$	$4.9 \cdot 10^{-4}$	$4.8 \cdot 10^{-5}$	$4.8 \cdot 10^{-6}$	$4.8 \cdot 10^{-7}$	$4.8 \cdot 10^{-8}$

Table 1: The maximal relative error from the specific linear equation solver. ϵ is $\log_{10}(|\frac{v-u}{u}|)$ where v is the calculated value and u is the closed-form solution.

4.4 Our method vs numpy.linalg.solve

Solving the linear equation systems using `numpy.linalg.solve` and our method for different amounts of data points results in:




Figure 4: Numpy solution vs my solution. They already seem similar.




Figure 5: Numpy solution vs my solution. It is becoming harder to tell them apart.

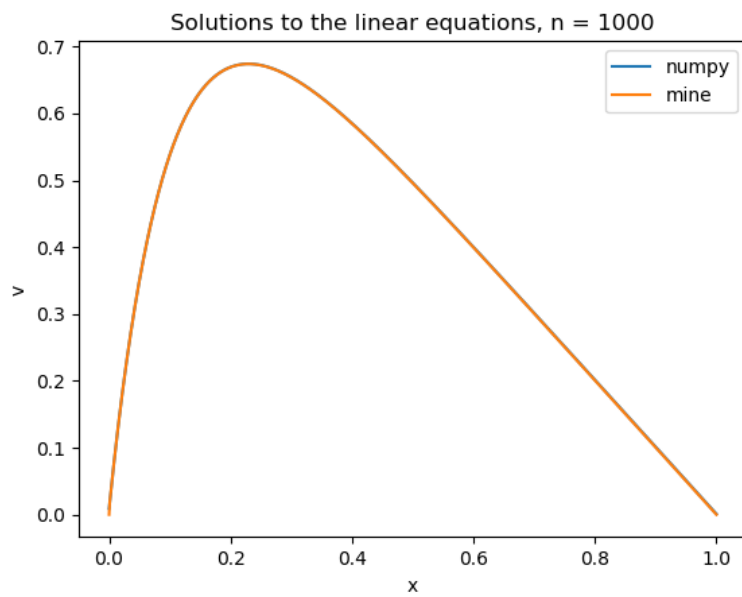


Figure 6: Numpy solution vs my solution. It is now impossible to tell them apart.

The elapsed time was:

Method	$n = 100$	$n = 1000$
numpy	0.001 s	0.015 s
mine	0.00097 s	0.010 s

Table 2: Computation time for the 2 methods.

5 Discussion

5.1 Row reduction methods.

The easiest, and least efficient, way to implement row reduction seems to be operating directly on a whole matrix using matrix notation. For this project, that was definitely not viable as we did not have enough memory for matrices larger than $10^4 \times 10^4$.

Removing all the zeros and operating on a $n \times 4$ matrix was also quite simple, and required a lot less computation time and memory. It is a shame work was not started on that immediately, as the memory problem was not realised until using $10^5 \times 10^5$ matrices was attempted and the first implementation took some time.

Making the specified linear equation solver was by far the biggest task this project. The immediate reaction was to find a way to omit any calculations that did not directly apply to \vec{f} . The implementation of this eluded success for quite a while, however. This resulted in the implementation of the method "slow_specific_linear_equation_solver", which follows the same principle as the method "linear_equation_solver". Later, implementation of the original idea was attempted once again, and succeeded.

The method "fast_specific_linear_equation_solver" seems to be around 5 times faster than the general one. This seems inconsistent with our calculation of FLOPS. This might be because we counted division to be 1 FLOP, while it is actually more costly than that.

5.2 Relative Error

From Table 1, it looks like solving the differential equation using row reduction is a first order method, causing the error to be reduced by a factor of 10 when the amount of data-points is increased by a factor of 10.

5.3 LU decomposition

We can see that the time elapsed inceases faster for the numpy method than for our method. This is no surprise, as the numpy method operates on the whole matrix. The numpy method assumedly has $\frac{2}{3}n^3$ FLOPS. It is

impossible to use it on a matrix as large as $10^5 \times 10^5$ as our computer lacks the memory.

6 Conclusion

From these experiments, we see the importance of making less generalized code when operating on very large matrices, as this can cut down on both memory usage and computation time.

7 Appendix

This is the code used in the project.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from time import time
4
5
6 class Project_solver():
7
8     def __init__(self, f, a = -1, b = 2, c = -1, M = True):
9         self.n = len(f)
10        self.h = 1/(self.n + 1)
11        self.f = f*self.h**2
12        self.a = a
13        self.b = b
14        self.c = c
15        if M == True:
16            self.builder()
17        else:
18            self.M = M
19    def builder(self):
20        """
21        Builds a matrix M
22        Used in the linear transformation
23        Mv = f
24        """
25        M = np.zeros((self.n, 4))
26        M[0, :] = np.array([self.b, self.c, 0, self.f[0]])
27        for i in range(1, self.n-1):
28            M[i, :] = np.array([self.a, self.b, self.c, self.f[i]])
29        M[-1, :] = np.array([0, self.a, self.b, self.f[-1]])
30        self.M = M
31        self.M[:, :-1] = self.M[:, :-1]
32
33
34    def linear_equation_solver(self):
35        """
36        Solves Mv = f for v
37        using forwards and backwards substitution
38        """
39        self.M[0, :] = self.M[0, :]/self.M[0, 0]
40        self.M[1, :] = self.M[1, :] - self.M[0, :]*self.M[1, 0]
41        for i in range(1, self.n-2):
42            self.M[i, 1:] = self.M[i, 1:]/self.M[i, 1]
43
44            multi_comp = self.M[i+1, 0]
45            for j in range(2):
46                self.M[i+1, j] = self.M[i+1, j] - self.M[i, j+1]*↔
47                    multi_comp
48            self.M[i+1, -1] = self.M[i+1, -1] - self.M[i, -1]*↔
49                    multi_comp
50        self.M[-2, :] = self.M[-2, :]/self.M[-2, 1]
51        self.M[-1, :] = self.M[-1, :] - self.M[-2, :]*self.M[-1, 1]
52        self.M[-1, :] = self.M[-1, :]/self.M[-1, -2]
53
54        self.M[-2, :] = self.M[-2, :] - self.M[-1, :]*self.M[-2, -2]
```

```

55     self.M[0, -1] = 0
56     self.thingyjing = self.M[:, -1]
57     for i in range(1, self.n-1):
58         multi_comp = self.M[-(i+1), -2]
59         self.M[-(i+1), -2] = \
60             self.M[-(i+1), -2] - self.M[-i, -3]*multi_comp
61         self.M[-(i+1), -1] = self.M[-(i+1), -1] - self.M[-i, -2]
62             -1]*multi_comp
63
64     self.M[0, :] = self.M[0, :] - self.M[1, :]*self.M[0, 1]
65     self.v = self.M[:, -1]
66     self.M[-1, -1] = 0
67     self.M[0, -1] = 0
68     self.v = self.M[:, -1]
69
70     def slow_specific_linear_equation_solver(self):
71         """
72         Solves  $Mv = f$  for  $v$ 
73         using the fact that the diagonal elements in  $M$  are all equal↔
74         .
75         More generalized.
76         """
77
78         self.M[0, :] = self.M[0, :]/self.M[0, 0]
79         self.M[1, :] = self.M[1, :] - self.M[0, :]*self.M[1, 0]
80
81         for i in range(1, self.n-2):
82             self.M[i, 2:] = self.M[i, 2:]/self.M[i, 1]
83             self.M[i+1, 1] = self.M[i+1, 1] + self.M[i, 2]
84             self.M[i+1, -1] = self.M[i+1, -1] + self.M[i, -1]
85
86         self.M[-2, :] = self.M[-2, :]/self.M[-2, 1]
87         self.M[-1, :] = self.M[-1, :] - self.M[-2, :]
88         self.M[-1, 2:] = self.M[-1, 2:]/self.M[-1, -2]
89         self.M[0, -1] = 0
90         self.M[-1, -1] = 0
91         for i in range(1, self.n-1):
92             self.M[-(i+1), -1] = self.M[-(i+1), -1] - \
93                 self.M[-i, -1]*self.M[-(i+1), -2]
94
95         self.M[-1, -1] = 0
96
97     def relative_error_finder(self):
98         """
99         finds the relative error epsilon
100         for the specific linear equation solver
101         epsilon = log10((v - u)/u)
102         """
103         i = np.arange(1, self.n+1)
104         numbers = (i)/(i+1)
105         self.fast_specific_linear_equation_solver(numbers)
106         x = np.linspace(0, 1, self.n)
107         u = (1 - (1 - np.exp(-10))*x - np.exp(-10*x))
108         v = self.v
109
110         self.epsilon = np.log10(np.abs((v[1:-1] - u[1:-1])/u[1:-1]))
111
112
113     def fast_specific_linear_equation_solver(self, numbers):
114         """

```

```

115         Solves  $Mv = f$  for  $v$ 
116         for the matrix
117          $\begin{bmatrix} 2 & -1 & \dots & f \\ -1 & 2 & -1 & \dots & f \\ \dots & -1 & 2 & -1 & \dots & f \\ \dots & 0 & 2 & -1 & \dots & f \end{bmatrix}$ 
118         using a list of numbers containing what to do with  $f$ .
119         """
120
121
122
123
124         v = self.f
125         v[0] = v[0]*numbers[0]
126         for i in range(1, self.n-1):
127             v[i] = (v[i] + v[i-1])*numbers[i]
128         v[-1] = 0
129         for i in range(1, self.n-1):
130             v[-(i+1)] = v[-(i+1)] + v[-i]*numbers[-(i+1)]
131         v[0] = 0
132         self.v = v
133     if __name__ == "__main__":
134         def trueplot(x):
135             return (1 - (1 - np.exp(-10))*x - np.exp(-10*x))
136
137         def linear_equation_tester():
138             n_list = [int(1e1), int(1e2), int(1e3)]
139             for n in n_list:
140                 x = np.linspace(0, 1, n)
141                 f = 100*np.exp(-10*x)
142                 objekt = Project_solver(f)
143                 objekt.linear_equation_solver()
144                 plt.plot(x, objekt.v)
145                 plt.plot(x, trueplot(x))
146                 plt.legend(["numerical", "closed-form solution"])
147                 plt.title("Numerical vs closed-form for n = %g" % n)
148                 plt.xlabel("x")
149                 plt.ylabel("u(x)")
150                 plt.show()
151
152         def computation_time_test():
153             n = int(1e6)
154             i = np.arange(1, n+1)
155             numbers = i/(i+1)
156
157             x = np.linspace(0, 1, n)
158             f = 100*np.exp(-10*x)
159
160             objekt = Project_solver(f)
161
162             start = time()
163             objekt.linear_equation_solver()
164             stop1 = time()
165             objekt.fast_specific_linear_equation_solver(numbers)
166             stop2 = time()
167
168             print("The first method took %g seconds" % (stop1 - start))
169             print("The second method took %g seconds" % (stop2 - stop1))
170             print((stop1 - start)/(stop2 - stop1))
171
172
173
174
175
176         def relative_error_tester():

```

```

177     ns = [10, 100, 1000, 10000, 100000, 1000000, 10000000]
178     error_max = []
179     for n in ns:
180         x = np.linspace(0, 1, n)
181         f = 100*np.exp(-10*x)
182
183         objekt = Project_solver(f)
184         objekt.relative_error_finder()
185         error_max.append(np.max(objekt.epsilon))
186     print(error_max)
187
188
189 def LU_tester():
190     ns = [10, 100, 1000]
191     a = -1
192     b = 2
193     c = -1
194
195     for n in ns:
196         x = np.linspace(0, 1, n)
197         f = 100*np.exp(-10*x)
198
199         M = np.zeros((n, n))
200         M[0, 0] = b
201         M[0, 1] = c
202         for i in range(1, n-1):
203             M[i, i-1] = a
204             M[i, i] = b
205             M[i, i+1] = c
206         M[-1, -2] = a
207         M[-1, -1] = b
208
209         h = 1/(n+1)
210         start = time()
211         v_np = np.linalg.solve(M, f*h**2)
212         stop1 = time()
213         objekt = Project_solver(f)
214         objekt.slow_specific_linear_equation_solver()
215         stop2 = time()
216         v_mine = objekt.M[:, -1]
217
218         plt.plot(x, v_np)
219         plt.plot(x, v_mine)
220         plt.legend(["numpy", "mine"])
221         plt.xlabel("x")
222         plt.ylabel("v")
223         plt.title("Solutions to the linear equations, n = %g" % n)
224         plt.show()
225         print(stop1 - start)
226         print(stop2 - stop1)
227     LU_tester()
228     #relative_error_tester()
229     #computation_time_test()

```